

Vikas Tadala

vikas32567@gmail.com | +91 9121032575 | India | [GitHub](#) | [LinkedIn](#)

SUMMARY

Technical Lead with 8+ years of experience building and scaling AI-driven, multi-tenant SaaS platforms. Specialized in LLM systems (RAG, agents, evaluation), distributed backend architecture, and cloud-native platforms serving 550K+ users. Proven track record of leading engineering teams in a fast-growing startup.

Experience

TECHNICAL LEAD | CHEZUBA CORP. | MAY 2021 – PRESENT

Leadership & Product:

- Worked directly under CEO, **Sukendra Reddy Rompally(Forbes 30 under 30)**
- Owned quarterly **Tech roadmaps**, prioritizing features and tech debt using ICE framework
- Participated in sales calls, **RFPs, and client onboarding** for 20+ enterprise customers
- Led and mentored **6-membered team** (Dev + QA)
- Established Agile sprint planning and delivery workflows using **Azure DevOps** (Boards, work items, dashboards), enabling predictable releases and improved team velocity
- Took care of **Tech resource planning** and interviewing candidates
- Handled the technical aspect of maintaining **SOC 2 Type 2, ISO 27001, 27017, 27018, GDPR, ROPA**

AI/LLM Systems:

- Built AI workflows using LangGraph tools and agents
- Designed RAG pipelines over datasets up to **8 million records**, optimizing embeddings and metadata
- Improved AI response accuracy from **~40% → ~90%**
- Reduced LLM latency from **90 seconds → 8 seconds**
- Reduced AI costs by 66% via token, prompt, and query optimization

Architecture & Scale:

- Designed and scaled a multi-tenant SaaS platform serving **400K+ enterprise users**
- Built event-driven services handling **200 TPS** with strict ordering guarantees
- Led system-level architectural decisions across databases, caching layers, message queues, and cloud infrastructure

Performance & Reliability:

- Offloaded long-running tasks to background workers, reducing API latency from **30s → 0.4s**
- Implemented request caching and queue-based retries to prevent dropped messages
- Achieved **99% uptime** across SaaS portals.
- Managed to maintain a **cache hit ratio of 90%** on the B2C portal
- Configured **application and resource-level monitors** with alerts for failure detection

APPLICATION DEVELOPER | CHEZUBA PVT. LTD. | SEPT 2018 – APRIL 2021

Application Development:

- Joined as founding engineer.
- Developed and maintained a B2C .NET application while maintaining coding standards
- Upgraded framework from **.NET Core 2.1 to .NET Core 3** and eventually to **.NET 6**
- Built a **chat solution from scratch** with access control that was linked to multiple clients
- Implemented **service level caching** using Azure Cache for Redis
- Reduced the **user drop-off rate from 70% to 35%** by improving page performance metrics
- Configured **CI/CD pipelines**, reducing deployment time from 2 hours to under 10 minutes
- Maintained a web portal built in Angular, which was exclusive to admin users
- Hosted **Apache SuperSet** for the business team to run analytics

SYSTEMS ENGINEER | INFOSYS | FEB 2017 – AUG 2018

Database maintenance:

- Build and maintain **ETL pipelines** for data migration
- Generate **reports in PowerBI** as per client needs

Education

INDIAN INSTITUTE OF TECHNOLOGY, DHANBAD | ELECTRICAL ENGINEERING | JULY 2012 – MAY 2016

Projects

AI RECOMMENDATION & AUDITING | PYTHON, LANGGRAPH, REACTJS

- Designed a LangGraph-based LLM workflow to **audit documents** and determine regulatory equivalency, **automating compliance** validation of NGOs at scale
- Built **conversational AI** supporting NGO discovery, CSR planning, and volunteering recommendations
- Maintained an NGO database of **8 million vectors** while maintaining low query response time
- Set up **logging using LangSmith** and analyse response accuracy
- Implemented **tools** for DB search, internet fetch, memory, and agent-based task routing
- Fine-tune Meta's Llama model to identify the scope of usability in the product
- Set up the Prefect3 engine to schedule the runs of Playwright jobs capable of scraping data

CORPORATE SaaS PORTALS | PYTHON, FRAPPE FRAMEWORK, FRAPPE CLOUD, STRIPE

- Building **highly customizable SaaS** solutions for diverse corporate clients with varying branding, workflow, data collection, and points crediting needs
- Implemented **SSO(OIDC)** and guided clients to enable federated identity on their cloud solution.
- Integrated **Stripe payment** solution for corporate donations processing
- Designed **real-time dashboards** for HR and CSR admins to track volunteer hours, donation totals, and qualitative social impact outcomes.
- Built reliable n8n workflows to migrate existing data from **Azure MS SQL to ClickUp**

TRANSACTION PROCESSING SERVICE | PYTHON, .NET, AWS LAMBDA, DYNAMODB, SQS

- Designed a high-throughput service **handling 200 TPS**
- Solved race conditions using transactional consistency in EF Core
- Leveraged Azure Table Storage partitioning for scalability
- Spawn and configure **AWS cloud services**: Lambda, SQS, DynamoDB, IAM

Personal Projects

SAPIENS | PYTHON, LANGCHAIN, FRAPPE FRAMEWORK

- A Frappe app that is capable of enabling RAG search on any DocType in the database
- Templates to generate embedding and search text from DB and user query, respectively

DELPHIS | EXPO REACT, NGINX, AZURE ENTRA ID

- An Android application capable of streaming to multiple platforms simultaneously
- Used **Azure Entra ID External Identities** as the identity provider
- Dynamically update **NGINX to restream** to social media platforms

Skills

- Languages: Python, C#, JavaScript
- Frameworks: FastAPI, LangChain, LangGraph, .NET Core, EF
- AI/ML Tools: Azure AI Foundry, Vertex AI, OpenAI ecosystem
- Cloud: Azure, AWS
- Databases: MariaDB, MS SQL, DynamoDB, Redis, Pinecone
- DevOps & Infra: Azure DevOps, GitHub Actions, Nginx, Docker
- Other: System Design, Architecture Reviews, Team Mentoring