

Vikas Tadala

vikas32567@gmail.com | +91 9121032575 | India | [GitHub](#) | [LinkedIn](#)

SUMMARY

Technical Lead with 8+ years of experience building and scaling AI-driven, multi-tenant SaaS platforms. Specialized in LLM systems (RAG, agents, evaluation), distributed backend architecture, and cloud-native platforms serving 550K+ users. Proven track record of leading engineering teams in a fast-growing startup.

Experience

TECHNICAL LEAD | CHEZUBA CORP. | MAY 2021 – DEC 2025

Leadership & Product:

- Owned quarterly **Tech roadmaps**, prioritizing features and tech debt using ICE framework
- Worked directly under **CEO, Sukendra Reddy Rompally(Forbes 30 under 30)**
- Participated in sales calls, **RFPs, and client onboarding** for 20+ enterprise customers
- Led and mentored **6-membered team** (Dev + QA)
- Established Agile sprint planning and delivery workflows using **Azure DevOps** (Boards, work items, dashboards), enabling predictable releases and improved team velocity
- Took care of **Tech resource planning** and interviewing candidates
- Handled the technical aspect of maintaining **SOC 2 Type 2, ISO 27001, 27017, 27018, GDPR, ROPA**

AI/LLM Systems:

- Built AI workflows using LangGraph tools and agents
- Designed RAG pipelines for datasets of over **8 million records**, optimizing embeddings and metadata
- Improved AI response accuracy from **~40% → ~90%**
- Reduced LLM latency from **90 seconds → 8 seconds**
- Reduced AI costs by 66% via token, prompt, and query optimization

Architecture & Scale:

- Designed and scaled a multi-tenant SaaS platform serving **400K+ enterprise users**
- Built event-driven services handling **200 TPS** with strict ordering guarantees
- Led system-level architectural decisions across databases, caching layers, message queues, and cloud infrastructure

Performance & Reliability:

- Offloaded long-running tasks to background workers, reducing API latency from **30s → 0.4s**
- Implemented request caching and queue-based retries to prevent dropped messages
- Achieved **99% uptime** across SaaS portals.
- Managed to maintain a **cache hit ratio of 90%** on the B2C portal
- Configured **application and resource-level monitors** with alerts for failure detection

APPLICATION DEVELOPER | CHEZUBA PVT. LTD. | SEPT 2018 – APRIL 2021

Application Development:

- Joined as founding engineer.
- Developed and maintained a B2C .NET application while maintaining coding standards
- Upgraded framework from **.NET Core 2.1 to .NET Core 3** and eventually to **.NET 6**
- Built a **chat solution from scratch** with access control that was linked to multiple clients
- Implemented **service level caching** using Azure Cache for Redis
- Reduced the **user drop-off rate from 70% to 35%** by improving page performance metrics
- Configured **CI/CD pipelines**, reducing deployment time from 2 hours to under 10 minutes
- Maintained a web portal built in Angular, which was exclusive to admin users
- Hosted **Apache SuperSet** for the business team to run analytics

SYSTEMS ENGINEER | INFOSYS | FEB 2017 – AUG 2018

Database maintenance:

- Build and maintain **ETL pipelines** for data migration
- Generate **reports in PowerBI** as per client needs

Education

INDIAN INSTITUTE OF TECHNOLOGY, DHANBAD | ELECTRICAL ENGINEERING | JULY 2012 – MAY 2016

Last three projects

EQUIVALENCY DETERMINATION | PYTHON, LANGGRAPH, AI FOUNDRY

- Built a LangGraph-based LLM workflow to automate NGO **document audits** and regulatory equivalency checks, with **dynamic planning** and document credibility validation
- Architected the base graph, planned agent/node responsibilities, and separated AI reasoning from python checks
- Improved planner-executor(nodes) accuracy via prompt tuning
- Integrated LangSmith for tracing, monitoring, debugging, and QA validation of output accuracy (including Kiln estimation)

CHEZUBA AI | PYTHON, LANGGRAPH, REACTJS, COSMOSDB

- Built a conversational AI platform for NGO discovery, CSR planning, and volunteering recommendations
- Designed flow for 8M+ vector NGO knowledge base with low-latency semantic search
- Led architecture decisions including LangGraph design, Azure AI Search POCs, and database selection
- Implemented LLM tools for database search and controlled internet retrieval, and fine-tuned Meta LLaMA model to define product usability boundaries
- Integrated LangSmith and orchestrated large-scale data collection using Prefect 3 with Playwright-based scraping jobs

RECOMMENDATION ENGINE | PYTHON, LANGCHAIN, PINECONE

- Built a RAG-based recommendation engine to personalize projects, fundraising campaigns, and volunteering opportunities
- Designed a table-agnostic RAG framework extensible across datasets
- Optimized prompts and retrieval for relevance and consistency
- Implemented combined location, skills/causes, and negation filters powering homepage recommendations

Personal Projects

RATTUS | PYTHON, LANGGRAPH, COSMOSDB

- A dual step python workflow that is capable answering user's analytical query from csv data
- Implemented schema inference and ingestion into vector and analytical databases
- Search workflow is capable of identifying relevant tables, reading schema, running queries, and answer user's query
- From learning planned a solution to handle 100GB+ csv data

DELPHIS | EXPO REACT, NGINX, AZURE ENTRA ID

- Developed an **Expo React Native** Android app for simultaneous multi-platform streaming
- Implemented **Azure Entra ID** External Identities for authentication
- Reconfigured **Nginx dynamically** to restream content

Skills

- Languages: Python, C#, JavaScript
- Frameworks: FastAPI, LangChain, LangGraph, .NET Core, EF
- AI/ML Tools: Azure AI Foundry, Vertex AI, OpenAI ecosystem
- Cloud: Azure, AWS
- Databases: MariaDB, MS SQL, DynamoDB, Redis, Pinecone
- DevOps & Infra: Azure DevOps, GitHub Actions, Nginx, Docker
- Other: System Design, Architecture Reviews, Team Mentoring